

Advanced Machine Learning - Project 1

Michał Gromadzki

March 2024

1 Methodology

This section describes methodology used in this project.

1.1 Datasets

Nine datasets have been selected from [UCI](#) and [OpenML](#). Three of those datasets have at most 10 variables and six have more than 10 variables. Figure 1 depicts distribution of datasets' dimensions.

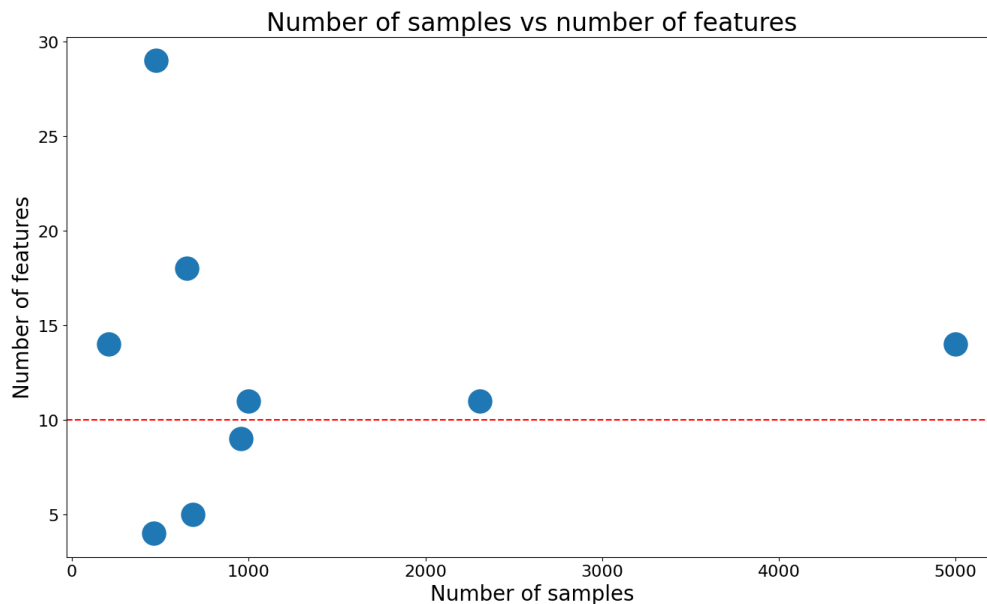


Figure 1: Number of samples and features

Datasets have at most 0.23% missing data.

In the next step, collinear variables have been removed from each dataset. Algorithm for removing collinear variables is described below.

Algorithm for removing collinear variables:

- Algorithm is iteration-based, in each iteration 1 variable can be removed
- In each iteration:
 1. Calculate variance inflation factor (VIF) for all variables
 2. Select all variables with VIF higher than $TH = 5$
 3. Delete variable with the highest VIF
- Algorithm stops when there is no variables with VIF over TH

1.2 Implementation

All optimization algorithms for parameter estimation in logistic regression have been implemented as classes in Python. All optimizers have the following attributes and methods. Attributes:

- `max_iter` (int) - Maximum number of iterations of the optimizer

- *tol* (float) - Minimum value of Frobenius norm of difference in parameters between iterations, used to determine convergence
- *coef_* (array) - Array containing parameters
- *interactions* (bool) - Whether to also include interactions of provided variables

Methods:

- *_add_interactions()* - Adds interactions to the provided variables
- *_sigmoid()* - Returns value of a sigmoid function for the provided input
- *fit()* - Trains the model, return history of training and number of iterations
- *predict_proba()* - Predicts probabilities based on the provided data
- *predict()* - Rounds the probabilities to class labels

Note that some implementations have individual attributes and methods necessary for their implementation.

1.3 Stopping rule

All algorithms have the same stopping rule to ensure fair comparison. At the end of each iteration of the algorithms the following value is computed:

$$diff = norm(coef - prev_coef)$$

where:

- *norm* - The Frobenius norm
- *coef* - Parameters in the current iteration
- *prev_coef* - Parameters in the previous iteration

If the *diff* is smaller than the pre-set *tol* value the training stops.

1.4 Training

All models have been trained on 9 different train-test splits. The *max_iter* and *tol* values have been set to 500 and 0.001, respectively. For SGD and Adam optimizer the *learning_rate* has been set to 0.01. Note that all variables have been scaled using standard scaler. To measure the performance of the models BalancedAccuracy was used on the test set.

2 Convergence analysis

To check for convergence, log-loss (negative log-likelihood) was computed on train set after each iteration. Figure 2 presents how the log-loss changes during training.

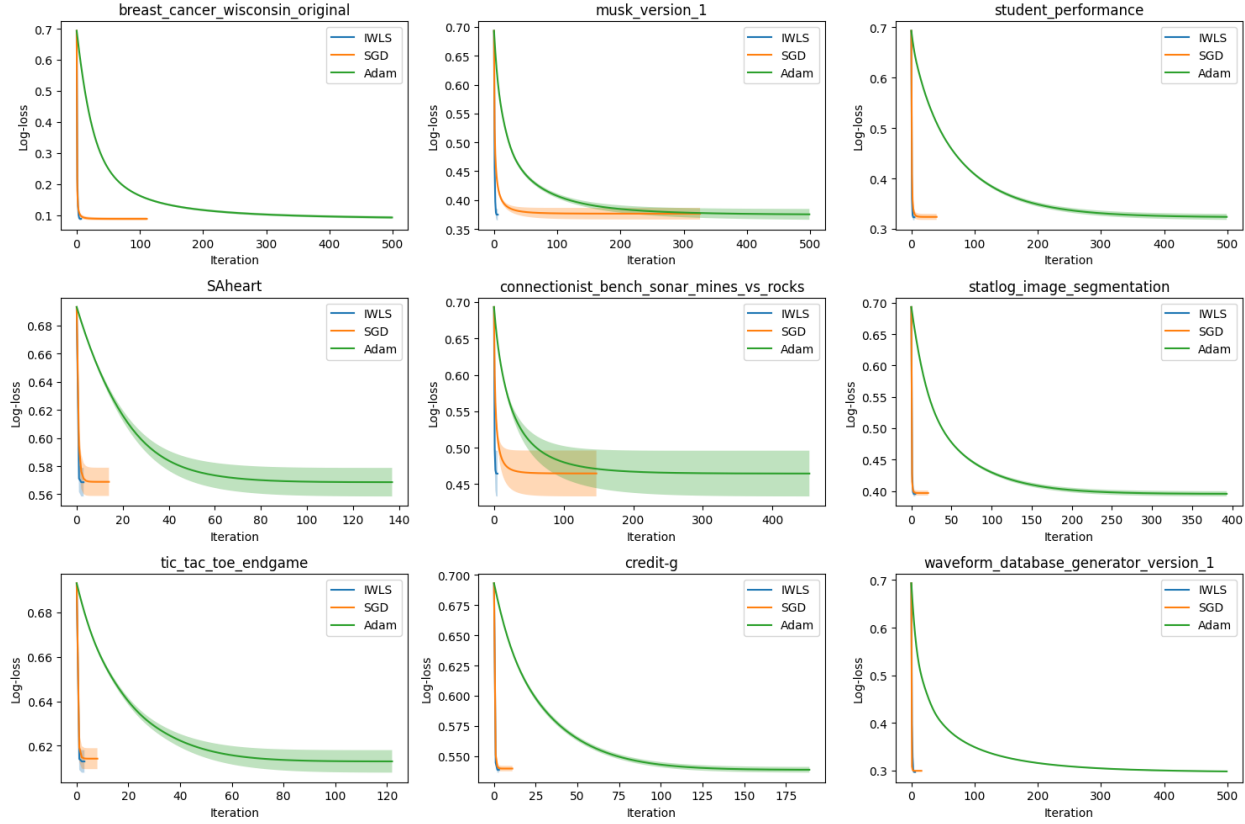


Figure 2: Log-loss value during training

Values at each iteration were computed as a mean of values from 9 different train-test splits. In all datasets final values of log-loss are similar for all methods, so we can expect similar classification performance from all methods. However, IWLS on average takes 5.7 iterations to finish training. On the other hand, SGD and Adam optimizers take much longer, 55 and 348 iterations, respectively.

3 Comparison of classification performance

Performance of the optimization algorithms has been tested on 9 datasets, along the implemented algorithms, LDA, QDA, DecisionTree and RandomForest were also used. Table 1 presents mean accuracy for each method.

Model	Accuracy	Std
IWLS	0.724	0.13
SGD	0.724	0.13
Adam	0.725	0.13
LDA	0.722	0.12
QDA	0.751	0.11
DT	0.748	0.14
RF	0.797	0.15

Table 1: Mean accuracy of each method

The best performing models is RandomForest. However if we focus only the implemented algorithms (IWLS, SGD, Adam), the best performing optimizer is Adam, by about 0.1%. Figure 3 depicts mean accuracy on each of the datasets.

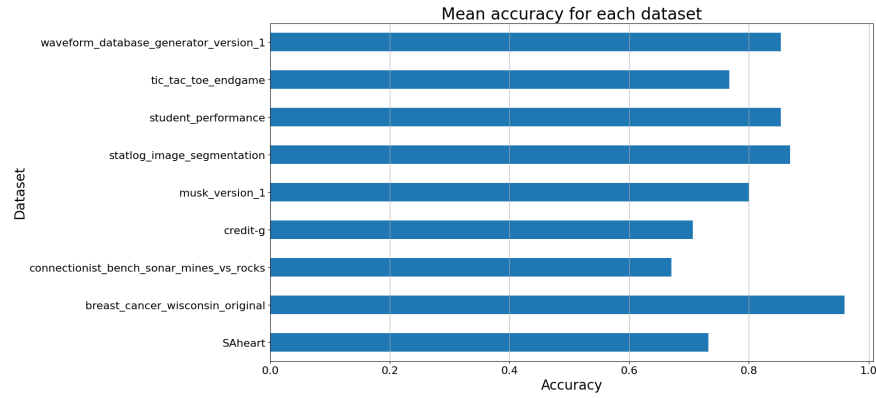


Figure 3: Mean accuracy on each of the datasets

Based on the Figure 3 we can tell which datasets are the easiest and which are the hardest. On average the hardest dataset is connectionist_bench_sonar_mines_vs_rocks and the easiest is breast_cancer_wisconsin_original. Figure 4 presents performance of all methods on all of the datasets.

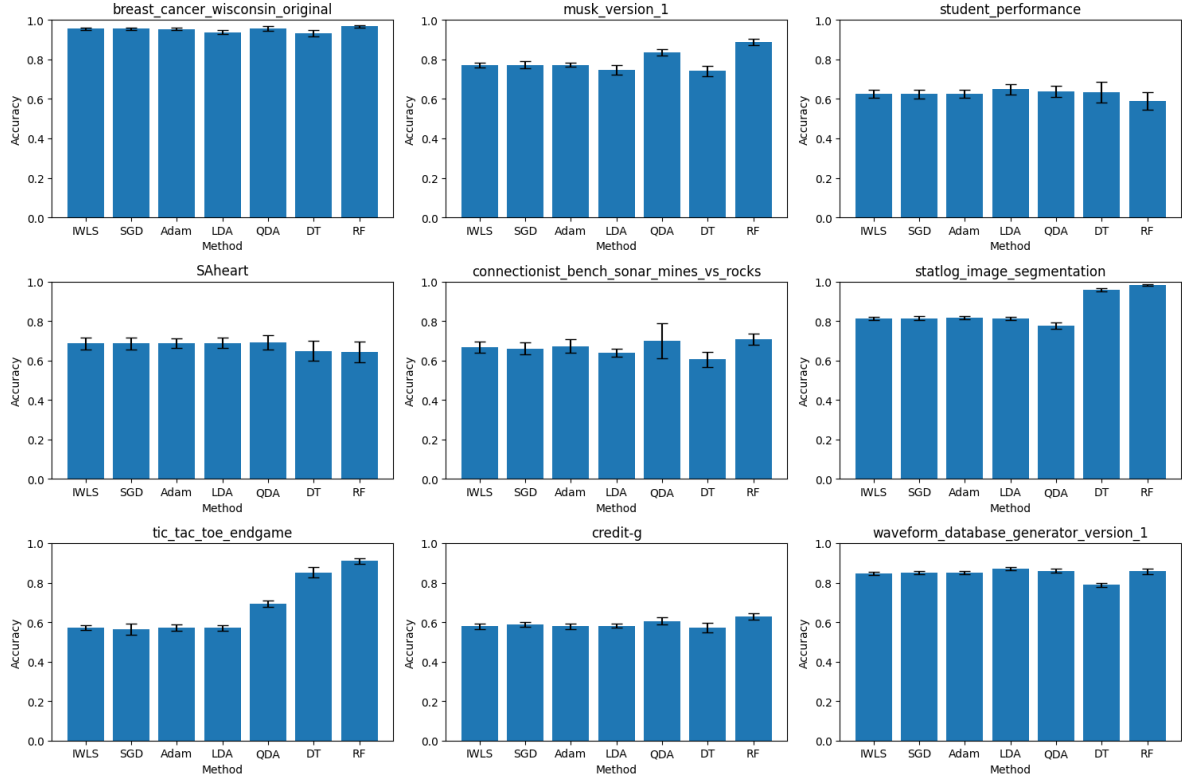


Figure 4: Accuracy of all methods on all datasets

RandomForest was the best performing algorithm in 8 out of 9 datasets. However if we only consider implemented algorithms, each one was the best performing one 3 times. In all cases the differences were small, the biggest one occurred in the connectionist_bench_sonar_mines_vs_rocks dataset and was equal to 1.2%. Moreover, one of the implemented algorithms was the worst only in one of the nine datasets.

4 Comparison of classification performance of models with and without interactions

For small datasets, which have less than 10 variables, two variations of logistic regression were tested. One with and one without interactions of the variables. Table 2 presents mean accuracy for each of the method.

Model	Accuracy	Std
IWLS	0.800	0.12
SGD	0.795	0.13
Adam	0.799	0.12
IWLS+INT	0.823	0.10
SGD+INT	0.825	0.10
Adam+INT	0.824	0.10

Table 2: Mean accuracy for each method

All models with interactions outperformed their counterparts without. The best performing one is SGD+INT. Using interactions amounted in about 2% increase in mean accuracy across all methods. Figure 5 presents accuracy of models with and without interactions on selected datasets.

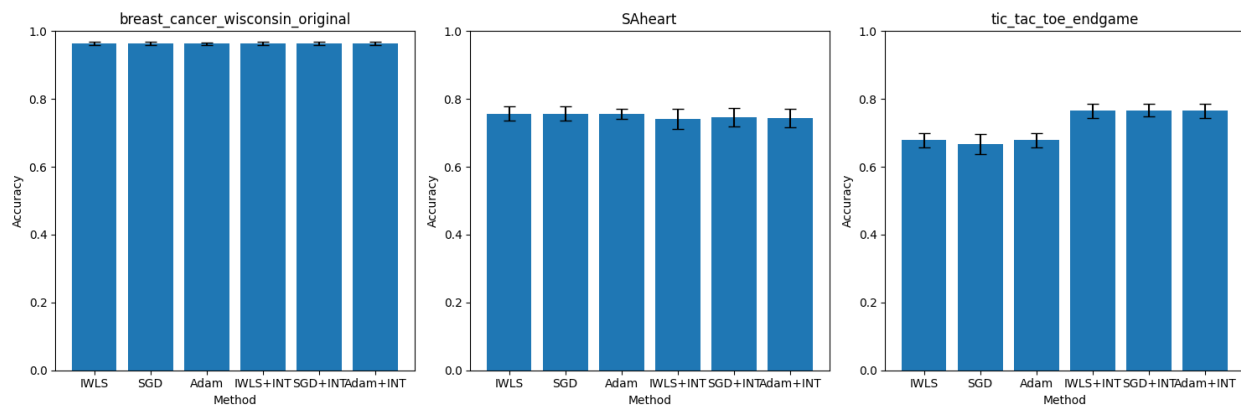


Figure 5: Accuracy of models with and without interactions

In two out of three datasets models with interactions outperformed those without. The biggest difference occurred in the tic_tac_toe_endgame dataset, where all models with interactions outperformed their counterparts without by about 10%.